

Home	Syllabus	What's next	Grades	Reading list	Handouts	Help	View your submissions for student1	View feedback / answers for student1	jmajikes
------	----------	-------------	--------	--------------	----------	------	------------------------------------	--------------------------------------	----------

Assessment was due by Sat, 2023-12-09 11:30:00

COMP 421 Final Exam Fall 2023

Note: There are a total of 110 points on this exams. The highest score you can get is a 100. Essentially there are 10 points that you can miss and still get a 100%.



Don't panic!

You have 180 minutes to finish the exam.

- You *must* stay in full screen mode. Points removed for leaving full screen mode
- You must hand this final exam in on time.
 - Points removed for late submissions.
 - Only your **first** submission will be accepted.
 - Avoid accidental submissions. Fill in your name when you are ready to submit.
- Points removed for accidental submissions.

I recommend that you have several pieces of scrap paper to doodle notes on during the exam.

Consider this final **closed book**.

You **MAY** use your hand written notes. They **MUST** be on paper as you may not switch screens after starting the exam.

You **MAY NOT** Google or use other external websites for **answers** or copy from a friend. Do not paste information into your exam unless it was copied from your exam. You **MAY NOT** receive help from anyone.

If you do not know the origin of material you should not paste it into this exam. All material pasted into this exam must originate from this exam. This

implies, but is not limited to, copying from previous assignments, copying from text messages, or copying from **any** website.

You **MUST** use the Google Chrome browser.

The browser will change input box color **green** to indicate correctness. A black or **red** box indicates an incorrect answer.

Note that HTML select statements with drop-downs are simple multiple choice questions. No highlighting of correct answers are done for select questions.

Green highlight should just assist you. If you believe your answer is correct and the input box did not turn **green**, continue on. Per the [syllabus](#), highlighting is simply an aide not a guarantee.

Note: For database queries that are applied to two databases, **two green lights** are required to get any credit for the question.

[SQL Tutorial Cheat Sheet](#)

Following are three SQL tutorial cheat sheets available from <http://www.sqltutorial.org>

SQL CHEAT SHEET <http://www.sqltutorial.org>

QUERYING DATA FROM A TABLE

SELECT c1, c2 FROM t;
Query data in columns c1, c2 from a table

SELECT * FROM t;
Query all rows and columns from a table

SELECT c1, c2 FROM t WHERE condition;
Query data and filter rows with a condition

SELECT DISTINCT c1 FROM t WHERE condition;
Query distinct rows from a table

SELECT c1, c2 FROM t ORDER BY c1 ASC [DESC];
Sort the result set in ascending or descending order

SELECT c1, c2 FROM t ORDER BY c1 LIMIT n OFFSET offset;
Skip *offset* of rows and return the next *n* rows

SELECT c1, aggregate(c2) FROM t GROUP BY c1;
Group rows using an aggregate function

SELECT c1, aggregate(c2) FROM t GROUP BY c1 HAVING condition;
Filter groups using HAVING clause

QUERYING FROM MULTIPLE TABLES

SELECT c1, c2 FROM t1 INNER JOIN t2 ON condition;
Inner join t1 and t2

SELECT c1, c2 FROM t1 LEFT JOIN t2 ON condition;
Left join t1 and t2

SELECT c1, c2 FROM t1 RIGHT JOIN t2 ON condition;
Right join t1 and t2

SELECT c1, c2 FROM t1 FULL OUTER JOIN t2 ON condition;
Perform full outer join

SELECT c1, c2 FROM t1 CROSS JOIN t2;
Produce a Cartesian product of rows in tables

SELECT c1, c2 FROM t1, t2;
Another way to perform cross join

SELECT c1, c2 FROM t1 A INNER JOIN t2 B ON condition;
Join t1 to itself using INNER JOIN clause

USING SQL OPERATORS

SELECT c1, c2 FROM t1 UNION [ALL] SELECT c1, c2 FROM t2;
Combine rows from two queries

SELECT c1, c2 FROM t1 INTERSECT SELECT c1, c2 FROM t2;
Return the intersection of two queries

SELECT c1, c2 FROM t1 MINUS SELECT c1, c2 FROM t2;
Subtract a result set from another result set

SELECT c1, c2 FROM t1 WHERE c1 [NOT] LIKE pattern;
Query rows using pattern matching % _

SELECT c1, c2 FROM t1 WHERE c1 [NOT] IN value_list;
Query rows in a list

SELECT c1, c2 FROM t1 WHERE c1 BETWEEN low AND high;
Query rows between two values

SELECT c1, c2 FROM t1 WHERE c1 IS [NOT] NULL;
Check if values in a table is NULL or not

SQL CHEAT SHEET <http://www.sqltutorial.org>

MANAGING TABLES

CREATE TABLE t (id INT PRIMARY KEY, name VARCHAR NOT NULL, price INT DEFAULT 0);
Create a new table with three columns

DROP TABLE t;
Delete the table from the database

ALTER TABLE t ADD column;
Add a new column to the table

ALTER TABLE t DROP COLUMN c;
Drop column c from the table

ALTER TABLE t ADD constraint;
Add a constraint

ALTER TABLE t DROP constraint;
Drop a constraint

ALTER TABLE t1 RENAME TO t2;
Rename a table from t1 to t2

ALTER TABLE t1 RENAME c1 TO c2;
Rename column c1 to c2

TRUNCATE TABLE t;
Remove all data in a table

USING SQL CONSTRAINTS

CREATE TABLE t(c1 INT, c2 INT, c3 VARCHAR, PRIMARY KEY (c1,c2));
Set c1 and c2 as a primary key

CREATE TABLE t1(c1 INT PRIMARY KEY, c2 INT, FOREIGN KEY (c2) REFERENCES t2(c2));
Set c2 column as a foreign key

CREATE TABLE t(c1 INT, c1 INT, UNIQUE(c2,c3));
Make the values in c1 and c2 unique

CREATE TABLE t(c1 INT, c2 INT, CHECK(c1 > 0 AND c1 >= c2));
Ensure c1 > 0 and values in c1 >= c2

CREATE TABLE t(c1 INT PRIMARY KEY, c2 VARCHAR NOT NULL);
Set values in c2 column not NULL

MODIFYING DATA

INSERT INTO t(column_list) VALUES(value_list);
Insert one row into a table

INSERT INTO t(column_list) VALUES (value_list), ..., (value_list), ...;
Insert multiple rows into a table

INSERT INTO t1(column_list) SELECT column_list FROM t2;
Insert rows from t2 into t1

UPDATE t SET c1 = new_value;
Update new value in the column c1 for all rows

UPDATE t SET c1 = new_value, c2 = new_value WHERE condition;
Update values in the column c1, c2 that match the condition

DELETE FROM t;
Delete all data in a table

DELETE FROM t WHERE condition;
Delete subset of rows in a table

SQL CHEAT SHEET <http://www.sqltutorial.org>

MANAGING VIEWS

CREATE VIEW `v(c1,c2)`
AS
SELECT `c1, c2`
FROM `t`;
 Create a new view that consists of c1 and c2

CREATE VIEW `v(c1,c2)`
AS
SELECT `c1, c2`
FROM `t`;
WITH [CASCADED | LOCAL] CHECK OPTION;
 Create a new view with check option

CREATE RECURSIVE VIEW `v`
AS
 select-statement -- anchor part
UNION [ALL]
 select-statement; -- recursive part
 Create a recursive view

CREATE TEMPORARY VIEW `v`
AS
SELECT `c1, c2`
FROM `t`;
 Create a temporary view

DROP VIEW `view_name`;
 Delete a view

MANAGING INDEXES

CREATE INDEX `idx_name`
ON `t(c1,c2)`;
 Create an index on c1 and c2 of the table t

CREATE UNIQUE INDEX `idx_name`
ON `t(c3,c4)`;
 Create a unique index on c3, c4 of the table t

DROP INDEX `idx_name`;
 Drop an index

SQL AGGREGATE FUNCTIONS

AVG returns the average of a list

COUNT returns the number of elements of a list

SUM returns the total of a list

MAX returns the maximum value in a list

MIN returns the minimum value in a list

MANAGING TRIGGERS

CREATE OR MODIFY TRIGGER `trigger_name`
WHEN EVENT
ON `table_name` **TRIGGER_TYPE**
EXECUTE `stored_procedure`;
 Create or modify a trigger

WHEN

- **BEFORE** – invoke before the event occurs
- **AFTER** – invoke after the event occurs

EVENT

- **INSERT** – invoke for INSERT
- **UPDATE** – invoke for UPDATE
- **DELETE** – invoke for DELETE

TRIGGER_TYPE

- **FOR EACH ROW**
- **FOR EACH STATEMENT**

CREATE TRIGGER `before_insert_person`
BEFORE INSERT
ON `person` **FOR EACH ROW**
EXECUTE `stored_procedure`;
 Create a trigger invoked before a new row is inserted into the person table

DROP TRIGGER `trigger_name`;
 Delete a specific trigger

Database Schema

Here are the tables you'll find for the database used in the midterm. Your queries will be run against two versions of the database. One of the databases will be much smaller and only contain a subset of the information.

```
CREATE TABLE classroom
(
    building      varchar(15),
    room_number   varchar(7),
    capacity      numeric(4,0),
    primary key (building, room_number)
)

CREATE TABLE department
(
    dept_name     varchar(20),
    building      varchar(15),
    budget        numeric(12,2) check (budget > 0),
    primary key (dept_name)
)

CREATE TABLE course
(
    course_id     varchar(8),
    title         varchar(50),
    dept_name     varchar(20),
    credits       numeric(2,0) check (credits > 0),
    primary key (course_id),
    foreign key (dept_name) references department (
```

```
        on delete set null
    )
CREATE INDEX idx_course_dept_name ON Course(dept_name)
CREATE TABLE instructor
    (ID          varchar(5), -- instructor's ID
     name        varchar(20) not null,
     dept_name   varchar(20),
     salary      numeric(8,2) check (salary > 29000)
     primary key (ID),
     foreign key (dept_name) references department (
         on delete set null
     )
CREATE INDEX idx_instructor_id ON Instructor(ID)
CREATE INDEX idx_instructor_dept_name ON Instructor(dept_name)
CREATE TABLE section
    (course_id   varchar(8),
     sec_id      varchar(8),
     semester    varchar(6)
     check (semester in ('Fall', 'Winter', 'Spring',
     year        numeric(4,0) check (year > 1701 and
     building     varchar(15),
     room_number  varchar(7),
     time_slot_id varchar(4),
     primary key (course_id, sec_id, semester, year)
     foreign key (course_id) references course (course_id)
     on delete cascade,
     foreign key (building, room_number) references
     on delete set null
     )
CREATE INDEX idx_section_year ON Section(year)
CREATE INDEX idx_section_semester ON Section(semester)
CREATE TABLE teaches
    (ID          varchar(5), -- instructor's ID
     course_id   varchar(8),
     sec_id      varchar(8),
     semester    varchar(6),
     year        numeric(4,0),
     primary key (ID, course_id, sec_id, semester, year)
     foreign key (course_id, sec_id, semester, year)
     references section (course_id, sec_id, semester, year)
```

```
        on delete cascade,
        foreign key (ID) references instructor (ID)
        on delete cascade
    )
CREATE INDEX idx_teaches_id ON Teaches(ID)
CREATE INDEX idx_teaches_year ON Teaches(year)
CREATE INDEX idx_teaches_semester ON Teaches(semester)
CREATE TABLE student
    (ID          varchar(5),
     name        varchar(20) not null,
     dept_name   varchar(20),
     tot_cred    numeric(3,0) check (tot_cred >= 0),
     primary key (ID),
     foreign key (dept_name) references department (
     on delete set null
    )
CREATE UNIQUE INDEX idx_student_id ON Student(ID)
CREATE INDEX idx_student_dept_name ON Student(dept_name)
CREATE TABLE takes
    (ID          varchar(5), -- Student ID
     course_id   varchar(8),
     sec_id      varchar(8),
     semester    varchar(6),
     year        numeric(4,0),
     grade       varchar(2),
     primary key (ID, course_id, sec_id, semester, year),
     foreign key (course_id, sec_id, semester, year)
     references section (course_id, sec_id, semester
     on delete cascade,
     foreign key (ID) references student (ID)
     on delete cascade
    )
CREATE INDEX idx_takes_id ON Takes(ID)
CREATE INDEX idx_takes_semester ON Takes(semester)
CREATE INDEX idx_takes_year ON Takes(year)
CREATE TABLE advisor
    (s_ID        varchar(5),
     i_ID        varchar(5),
     primary key (s_ID),
     foreign key (i_ID) references instructor (ID)
```

```

on delete set null,
foreign key (s_ID) references student (ID)
on delete cascade
)

```

```

CREATE INDEX idx_advisor_instructor_id ON Advisor(i_
CREATE TABLE time_slot

```

```

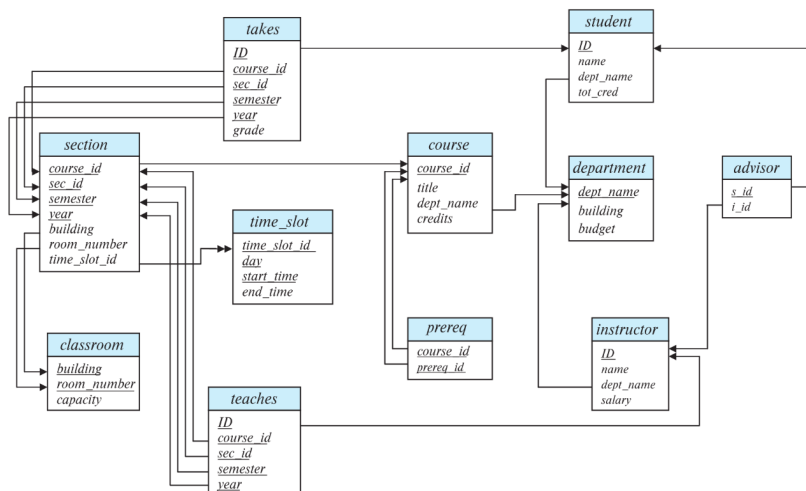
(time_slot_id   varchar(4),
day             varchar(1),
start_hr        numeric(2) check (start_hr >= 0
start_min       numeric(2) check (start_min >= 0
end_hr          numeric(2) check (end_hr >= 0 an
end_min         numeric(2) check (end_min >= 0 a
primary key (time_slot_id, day, start_hr, start
)

```

```

CREATE TABLE prereq
(course_id       varchar(8),
prereq_id       varchar(8),
primary key (course_id, prereq_id),
foreign key (course_id) references course (cour
on delete cascade,
foreign key (prereq_id) references course (cour
)

```



Scratch area

The following scratch space can be used to help develop and test queries against a database described above. The database used by the exam grader will be different.

Execute Minimize Output

Questions For a total of 110 points

SQL Queries 60 points

In this section, you will write SQL queries for the university described in the book. Your queries will be tested immediately against two different databases. If your queries output matches the expected output, the displayed answers will be outlined in green. Your actual score will be determined when your query is tested against a different database but green feedback should mean that you are on track to receive full credit.

List.Instructors.1: List the names of the instructors in the department of Accounting.

List the names in reverse alphabetical order.

Execute 5 points Minimize Output

List.Students.And.Instructors.1: List the names of people in the department of Accounting. This should be a single column table with the names of both the instructors and the students in the department of Accounting.

List the names in reverse alphabetical order.

Execute 7.5 points Minimize Output

List.Instructor.1.3: List the names of the instructors in the department of Accounting who have a name with exactly 3 letters in their name. List the names in alphabetical order.

Execute

5 points

Minimize Output

List.Instructor.Did.Not.Teach.All.1: List the names of the instructors in the department of Accounting who have **never** taught at least one Accounting course.

Another way to say this is: list all the names of Accounting instructors except those who taught all of the Accounting courses.

For example, if the Accounting department offers twelve courses, your answer should include any Accounting instructor's name who hasn't taught at least one section of every offered Accounting course.

Execute

7.5 points

Minimize Output

Update.Credits.1: Update the credits for Accounting courses such that they are worth 50% more than before.

NOTE₁: With all SQL statements like this that modify the database, you either need to refresh the web page before rerunning or handle the fact that the statement was previously run.

NOTE₂: The grader will run your modifications against both databases which should report **two** empty green boxes. Then the grader will run two select statments against both databases to verify the modification which should be in **four** more green boxes. A correct answer should have a total of **six** green boxes

Execute

7.5 points

Minimize Output

Relational.Algebra.1: Create a SQL query for the relational expression:

$$\Pi_{\text{title}}(\sigma_{\text{semester}='Fall' \wedge \text{dept_name}='Accounting'}(\rho_s(\text{Section}) \bowtie_{\text{c.course_id}=\text{s.course_id}} \rho_c(\text{Course})))$$

—

$$\Pi_{\text{title}}(\sigma_{\text{semester}='Winter' \wedge \text{dept_name}='Accounting'}(\rho_s(\text{Section}) \bowtie_{\text{c.course_id}=\text{s.course_id}} \rho_c(\text{Course})))$$

—

$$\Pi_{\text{title}}(\sigma_{\text{semester}='Spring' \wedge \text{dept_name}='Accounting'}(\rho_s(\text{Section}) \bowtie_{\text{c.course_id}=\text{s.course_id}} \rho_c(\text{Course})))$$

—

$$\Pi_{\text{title}}(\sigma_{\text{semester}='Summer' \wedge \text{dept_name}='Accounting'}(\rho_s(\text{Section}) \bowtie_{\text{c.course_id}=\text{s.course_id}} \rho_c(\text{Course})))$$

Execute

5 points

Minimize Output

Largest.Classroom.1: List the building and room_number of the classroom(s) of a Accounting class section that has the largest capacity.

Execute 7.5 points Minimize Output

List.Department.Budget.Maximum.Salary.1: List all the department names, their budgets, and the maximum instructor salary for each department. List the departments by largest budget to smallest budget.
NOTE: You can be assured that all departments have at least one instructor.

Execute 5 points Minimize Output

List.Students.With.Their.Advisor.1: List the name of each student in the Accounting department who has an advisor. In addition to the name of the

student, give the name of the student's advisor.

Note: Only the student has to be in the Accounting department.

Execute 5 points Minimize Output

List.Students.And.Advisor.If.Exists.1: List the name of each student in the Accounting department and if they have an advisor list the advisor name. If the student doesn't have an advisor leave the advisor name NULL.

Execute 5 points Minimize Output

Query processing 20 points

In this section, you will show your knowledge of Chapter 15 query processing

Calculator: You may use this box as a calculator. Just type in any expression that can be evaluated by a JavaScript eval. For example, $2 ** 0.5$ will show you the result of the square root of 2 in the red box to the right. Leaving the box empty or filled in will **NOT** affect your grade.

undefined

Let relation r_1 have 20,000 tuples where each block holds 25 tuples. Let relation r_2 have 45,000 tuples where each block holds 30 tuples.

Nested.Loop.R1.Outer.1: How many block transfers are required if a nested-loop join is done with r_1 as the outer loop and there is only one buffer for r_1 and one buffer for r_2 ?

NOTE: As previously done in class, do not count the number of block transfers required for output of the nested-loop join.

2 points

Nested.Loop.R2.Outer.1: How many block transfers are required if a nested-loop join is done with r_2 as the outer loop and there is only one buffer for r_1 and one buffer for r_2 ?

NOTE: As previously done in class, do not count the number of block transfers required for output of the nested-loop join.

2 points

Nested.Loop.Which.Outer.1: In order to minimized the number of block transfers, which relation should be used for the outer loop?

 1 point

Block.Nested.Loop.R1.Outer.1: How many block transfers are required if a block nested-loop join is used with r_1 as the outer relation?

Assume the enhanced version of block nested-loop join where M blocks are used for buffering of input of r_1 and r_2 and one buffer reserved to collect the output of the join.

NOTE₁: As previously done in class, do not count the number of block transfers required for output of the block nested-loop join.

NOTE₂: Remember, the evaluation box allows you to use functions floor(), ceil(), and log().

5 points

Sort.Merge.R1.1: Assuming there are 25 blocks of memory available for sorting r_1 , complete the following table below. For each pass, how many runs exists in that pass and the number of blocks in the largest run.

NOTE₁: The table has notation for up to four passes. If the merge-sort completes in less than four passes, those *extra* passes would note that one run exists and that the largest run is the number of blocks in r_1 .

NOTE₂: Remember, the evaluation box allows you to use functions floor(), ceil(), and log().

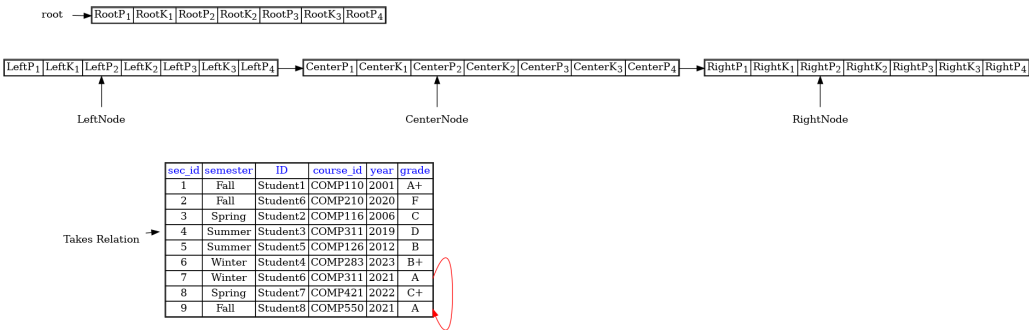
Pass number	Number runs that exist	Number of blocks in the largest run
Pass 1	<input type="text"/>	<input type="text"/>
Pass 2	<input type="text"/>	<input type="text"/>
Pass 3	<input type="text"/>	<input type="text"/>
Pass 4	<input type="text"/>	<input type="text"/>

10 points

B-plus trees 30 points

In this section, you will show your knowledge of B^+ trees.

Assume the following B^+ tree with $n = 4$ and the **Takes** relation and a search key of grade:



Complete.B.Plus.Tree.1: Complete the indices for relation Takes using a search key of grade. You should assume the same notation used in class and the same algorithm used in the book.

To facilitate having only one single correct answer, field CenterK3 must be 'C+'. It may help you to start building your tree from field CenterK3.


NOTE: Ascending alphanumeric sort order for grades is A, A+, A-, B, B+, B-, C, C+, C-, D, D+, D-, F

Field	Value
RootP1	<input type="text"/>

RootK1	<input type="text"/>
RootP2	<input type="text"/>
RootK2	<input type="text"/>
RootP3	<input type="text"/>
RootK3	<input type="text"/>
RootP4	<input type="text"/>
LeftP1	<input type="text"/>
LeftK1	<input type="text"/>
LeftP2	<input type="text"/>
LeftK2	<input type="text"/>
LeftP3	<input type="text"/>
LeftK3	<input type="text"/>
LeftP4	<input type="text"/>
CenterP1	<input type="text"/>
CenterK1	<input type="text"/>
CenterP2	<input type="text"/>
CenterK2	<input type="text"/>
CenterP3	<input type="text"/>
CenterK3	<input type="text"/>
CenterP4	<input type="text"/>
RightP1	<input type="text"/>
RightK1	<input type="text"/>
RightP2	<input type="text"/>
RightK2	<input type="text"/>
RightP3	<input type="text"/>
RightK3	<input type="text"/>
RightP4	<input type="text"/>

15 points

Assume the following B^+ tree with $n = 3$ and the **Takes** relation and a search key of grade:




Takes Relation →

sec_id	semester	ID	course_id	year	grade
3	Spring	Student2	COMP116	2006	C
4	Summer	Student3	COMP311	2019	D
5	Summer	Student5	COMP126	2012	B
7	Winter	Student6	COMP311	2021	A

Insert.B.Plus.Tree.1: Complete the following table of grade search keys and pointers after an insert of record '8 Spring Student7 COMP421 2022 C+'.

Note: You can be assured that RootP1 points to LeftNode, RootP2 points to CenterNode, and RootP3 points to RightNode. **NOTE:** Ascending alphanumeric sort order for grades is A, A+, A-, B, B+, B-, C, C+, C-, D, D+, D-, F



Takes Relation →

sec_id	semester	ID	course_id	year	grade
3	Spring	Student2	COMP116	2006	C
4	Summer	Student3	COMP311	2019	D
5	Summer	Student5	COMP126	2012	B
7	Winter	Student6	COMP311	2021	A
8	Spring	Student7	COMP421	2022	C+

Field	Value
RootP1	<input type="text"/>
RootK1	<input type="text"/>
RootP2	<input type="text"/>
RootK2	<input type="text"/>
RootP3	<input type="text"/>
LeftP1	<input type="text"/>
LeftK1	<input type="text"/>
LeftP2	<input type="text"/>
LeftK2	<input type="text"/>
LeftP3	<input type="text"/>
CenterP1	<input type="text"/>
CenterK1	<input type="text"/>
CenterP2	<input type="text"/>

CenterK2	<input type="text"/>
CenterP3	<input type="text"/>
RightP1	<input type="text"/>
RightK1	<input type="text"/>
RightP2	<input type="text"/>
RightK2	<input type="text"/>
RightP3	<input type="text"/>

15 points
