

Home	Syllabus	What's next	Grades	Reading list	Handouts	Help	View your submissions for student1	View feedback / answers for student1	jmajikes
------	----------	-------------	--------	--------------	----------	------	------------------------------------	--------------------------------------	----------

Assessment due **before** Fri, 2024-05-31 00:00:00.

COMP 421 Midterm 1 Fall 2023

Note: There are a total of 100 points on this exams.

Don't panic!

You have 75 minutes to finish the exam.

- You *should* stay in full screen mode.
 - A <ctrl-f> find will give you a warning, which is OK
- You must hand this midterm during class or prearranged midterm timeframe.
 - Avoid accidental submissions. Fill in your name when you are ready to submit.
 - Only your **first** submission will be accepted/graded for full credit.
 - If and only if you submit your exam during class you will have a chance to resubmit it for reduced credit. See [subsequent submissions on the syllabus for more information](#).
 - After all students have submitted and the grader has finished, the submit button will be enabled.
 - You have until 2023-10-17 11:59:59 to submit any subsequent submissions
 - Plan your time judiciously!

I recommend that you have several pieces of scrap paper to doodle notes on during the exam. I *strongly* recommend you read the whole exam and begin with questions you know how to solve quickly. Some questions will be harder or take longer than others; don't spend all your time on one question worth only a few points!



Consider this midterm **closed book**.

You can **NOT** reference other online homeworks, worksheets, etc.

You can use your notes or other things printed out. They should be on paper as you may not switch screens after starting the exam.

You **MAY NOT** Google for anything, You **MAY NOT** leave this website, you **MAY NOT** visit any websites, and you **MAY NOT** copy from a friend. Do not paste information into your midterm unless you know it came from your midterm. You **MAY NOT** receive help from anyone.

If you do not know the origin of material you should not paste it into this exam. All material pasted into this exam must originate from this exam. This implies, but is not limited to, copying from previous assignments, copying from text messages, or copying from **any** website.

You **MUST** use the Google Chrome browser.

The instruction team will **not** answer questions about course content, SQL syntax, etc. We will only deal with issues related to exam implementation.

If your browser hangs, for example because of a bad SQL query, simply kill the page and refresh. It *should* restore all your work even if it doesn't re-evaluate all answers, color-highlight boxes, etc.

You may **NOT** leave the classroom before you submit your exam. When you submit your exam you **must enter the code** displayed on the screen at the front of the class or given to you by ARS. `

You **must not** use your computer or phone in the classroom after you submit your exam. After submitting your exam, simply leave the classroom or ARS.

The browser will change input box color **green** to indicate correctness. A black or **red** box indicates an incorrect answer.

Note that HTML select statements with drop-downs are simple multiple choice questions. No highlighting of correct answers are done for select questions.

Green highlight should just assist you. If you believe your answer is correct and the input box did not turn **green**, continue on. Per the [syllabus](#), highlighting is simply an aide not a guarantee.

Note: For database queries that are applied to two databases, **two green lights** are required to get any credit for the question.

[SQL Tutorial Cheat Sheet](#)

Following are three SQL tutorial cheat sheets available from <http://www.sqltutorial.org>

SQL CHEAT SHEET <http://www.sqltutorial.org>

QUERYING DATA FROM A TABLE

SELECT c1, c2 FROM t;
Query data in columns c1, c2 from a table

SELECT * FROM t;
Query all rows and columns from a table

**SELECT c1, c2 FROM t
WHERE condition;**
Query data and filter rows with a condition

**SELECT DISTINCT c1 FROM t
WHERE condition;**
Query distinct rows from a table

**SELECT c1, c2 FROM t
ORDER BY c1 ASC [DESC];**
Sort the result set in ascending or descending order

**SELECT c1, c2 FROM t
ORDER BY c1
LIMIT n OFFSET offset;**
Skip offset of rows and return the next n rows

**SELECT c1, aggregate(c2)
FROM t
GROUP BY c1;**
Group rows using an aggregate function

**SELECT c1, aggregate(c2)
FROM t
GROUP BY c1
HAVING condition;**
Filter groups using HAVING clause

QUERYING FROM MULTIPLE TABLES

**SELECT c1, c2
FROM t1
INNER JOIN t2 ON condition;**
Inner join t1 and t2

**SELECT c1, c2
FROM t1
LEFT JOIN t2 ON condition;**
Left join t1 and t2

**SELECT c1, c2
FROM t1
RIGHT JOIN t2 ON condition;**
Right join t1 and t2

**SELECT c1, c2
FROM t1
FULL OUTER JOIN t2 ON condition;**
Perform full outer join

**SELECT c1, c2
FROM t1
CROSS JOIN t2;**
Produce a Cartesian product of rows in tables

**SELECT c1, c2
FROM t1, t2;**
Another way to perform cross join

**SELECT c1, c2
FROM t1 A
INNER JOIN t2 B ON condition;**
Join t1 to itself using INNER JOIN clause

USING SQL OPERATORS

**SELECT c1, c2 FROM t1
UNION [ALL]
SELECT c1, c2 FROM t2;**
Combine rows from two queries

**SELECT c1, c2 FROM t1
INTERSECT
SELECT c1, c2 FROM t2;**
Return the intersection of two queries

**SELECT c1, c2 FROM t1
MINUS
SELECT c1, c2 FROM t2;**
Subtract a result set from another result set

**SELECT c1, c2 FROM t1
WHERE c1 [NOT] LIKE pattern;**
Query rows using pattern matching % _

**SELECT c1, c2 FROM t
WHERE c1 [NOT] IN value_list;**
Query rows in a list

**SELECT c1, c2 FROM t
WHERE c1 BETWEEN low AND high;**
Query rows between two values

**SELECT c1, c2 FROM t
WHERE c1 IS [NOT] NULL;**
Check if values in a table is NULL or not

SQL CHEAT SHEET <http://www.sqltutorial.org>

MANAGING TABLES

**CREATE TABLE t (
id INT PRIMARY KEY,
name VARCHAR NOT NULL,
price INT DEFAULT 0
);**
Create a new table with three columns

DROP TABLE t;
Delete the table from the database

ALTER TABLE t ADD column;
Add a new column to the table

ALTER TABLE t DROP COLUMN c;
Drop column c from the table

ALTER TABLE t ADD constraint;
Add a constraint

ALTER TABLE t DROP constraint;
Drop a constraint

ALTER TABLE t1 RENAME TO t2;
Rename a table from t1 to t2

ALTER TABLE t1 RENAME c1 TO c2;
Rename column c1 to c2

TRUNCATE TABLE t;
Remove all data in a table

USING SQL CONSTRAINTS

**CREATE TABLE t(
c1 INT, c2 INT, c3 VARCHAR,
PRIMARY KEY (c1,c2)
);**
Set c1 and c2 as a primary key

**CREATE TABLE t1(
c1 INT PRIMARY KEY,
c2 INT,
FOREIGN KEY (c2) REFERENCES t2(c2)
);**
Set c2 column as a foreign key

**CREATE TABLE t(
c1 INT, c1 INT,
UNIQUE(c2,c3)
);**
Make the values in c1 and c2 unique

**CREATE TABLE t(
c1 INT, c2 INT,
CHECK(c1 > 0 AND c1 >= c2)
);**
Ensure c1 > 0 and values in c1 >= c2

**CREATE TABLE t(
c1 INT PRIMARY KEY,
c2 VARCHAR NOT NULL
);**
Set values in c2 column not NULL

MODIFYING DATA

**INSERT INTO t(column_list)
VALUES(value_list);**
Insert one row into a table

**INSERT INTO t(column_list)
VALUES (value_list), ...;**
Insert multiple rows into a table

**INSERT INTO t1(column_list)
SELECT column_list
FROM t2;**
Insert rows from t2 into t1

**UPDATE t
SET c1 = new_value;**
Update new value in the column c1 for all rows

**UPDATE t
SET c1 = new_value,
c2 = new_value
WHERE condition;**
Update values in the column c1, c2 that match the condition

DELETE FROM t;
Delete all data in a table

**DELETE FROM t
WHERE condition;**
Delete subset of rows in a table

SQL CHEAT SHEET <http://www.sqltutorial.org>

MANAGING VIEWS

CREATE VIEW `v(c1,c2)`
AS
SELECT `c1, c2`
FROM `t`;
 Create a new view that consists of c1 and c2

CREATE VIEW `v(c1,c2)`
AS
SELECT `c1, c2`
FROM `t`;
WITH [CASCADED | LOCAL] CHECK OPTION;
 Create a new view with check option

CREATE RECURSIVE VIEW `v`
AS
 select-statement; -- anchor part
UNION [ALL]
 select-statement; -- recursive part
 Create a recursive view

CREATE TEMPORARY VIEW `v`
AS
SELECT `c1, c2`
FROM `t`;
 Create a temporary view

DROP VIEW `view_name;`
 Delete a view

MANAGING INDEXES

CREATE INDEX `idx_name`
ON `t(c1,c2);`
 Create an index on c1 and c2 of the table t

CREATE UNIQUE INDEX `idx_name`
ON `t(c3,c4);`
 Create a unique index on c3, c4 of the table t

DROP INDEX `idx_name;`
 Drop an index

SQL AGGREGATE FUNCTIONS

AVG returns the average of a list

COUNT returns the number of elements of a list

SUM returns the total of a list

MAX returns the maximum value in a list

MIN returns the minimum value in a list

MANAGING TRIGGERS

CREATE OR MODIFY TRIGGER `trigger_name`
WHEN EVENT
ON `table_name` **TRIGGER_TYPE**
EXECUTE `stored_procedure;`
 Create or modify a trigger

WHEN

- BEFORE** – invoke before the event occurs
- AFTER** – invoke after the event occurs

EVENT

- INSERT** – invoke for INSERT
- UPDATE** – invoke for UPDATE
- DELETE** – invoke for DELETE

TRIGGER_TYPE

- FOR EACH ROW**
- FOR EACH STATEMENT**

CREATE TRIGGER `before_insert_person`
BEFORE INSERT
ON `person` **FOR EACH ROW**
EXECUTE `stored_procedure;`
 Create a trigger invoked before a new row is inserted into the person table

DROP TRIGGER `trigger_name;`
 Delete a specific trigger

Following is a SQL tutorial cheat sheets available from <http://learnsql.com>

SQL Cheat Sheet <https://learnsql.com/blog/sql-for-data-analysis-cheat-sheet/>

SQL

SQL or Structured Query Language, is a language for talking to databases. It lets you select specific data and build complex reports. Today, SQL is a universal language of data, used in practically all technologies that process data.

SELECT

Fetch the id and name columns from the product table:
SELECT `id, name`
FROM `product;`

Concatenate the name and the description to fetch the full description of the products:
SELECT `name || ' - ' || description`
FROM `product;`

Fetch names of products with prices above 15:
SELECT `name`
FROM `product`
WHERE `price > 15;`

Fetch names of products with prices between 50 and 150:
SELECT `name`
FROM `product`
WHERE `price BETWEEN 50 AND 150;`

Fetch names of products that are not watches:
SELECT `name`
FROM `product`
WHERE `name != 'watch';`

Fetch names of products that start with a 'P' or end with an 's':
SELECT `name`
FROM `product`
WHERE `name LIKE 'P%' OR name LIKE '%s';`

Fetch names of products that start with any letter followed by "rain" (like "Train" or "grain"):
SELECT `name`
FROM `product`
WHERE `name LIKE '%_rain';`

Fetch names of products with non-null prices:
SELECT `name`
FROM `product`
WHERE `price IS NOT NULL;`

GROUP BY

name	category	category	count
Kettle	Kitchen	Kitchen	3
Mixer	Kitchen	Kitchen	3
Jeans	Clothing	Clothing	3
Sneakers	Clothing	Electronics	2
Jeggings	Clothing		
Smart TV	Electronics		
Laptop	Electronics		

AGGREGATE FUNCTIONS

Count the number of products:
SELECT `COUNT(*)`
FROM `product;`

Count the number of products with non-null prices:
SELECT `COUNT(price)`
FROM `product;`

Count the number of unique category values:
SELECT `COUNT(DISTINCT category)`
FROM `product;`

Get the lowest and the highest product price:
SELECT `MIN(price), MAX(price)`
FROM `product;`

Find the total price of products for each category:
SELECT `category, SUM(price)`
FROM `product`
GROUP BY `category;`

Find the average price of products for each category whose average is above 3.0:
SELECT `category, AVG(price)`
FROM `product`
GROUP BY `category`
HAVING `AVG(price) > 3.0;`

ORDER BY

Fetch product names sorted by the price column in the default Ascending order:
SELECT `name`
FROM `product`
ORDER BY `price [ASC];`

Fetch product names sorted by the price column in Descending order:
SELECT `name`
FROM `product`
ORDER BY `price DESC;`

COMPUTATIONS

Use `+`, `-`, `*`, `/` to do basic math. To get the number of seconds in a week:
SELECT `60 * 60 * 24 * 7;`
-- result: 604800

ROUNDING NUMBERS

Round a number to its nearest integer:
SELECT `ROUND(1234.56789);`
-- result: 1235

Round a number to two decimal places:
SELECT `ROUND(AVG(price), 2)`
FROM `product`
WHERE `category_id = 21;`
-- result: 124.50

TROUBLESHOOTING

INTEGER DIVISION

In PostgreSQL and SQL Server, the `/` operator performs integer division for integer arguments. If you do not see the number of decimal places you expect, it is because you are dividing between two integers. Cast one to decimal:
123 / 2 -- result: 61
CAST(123 AS decimal) / 2 -- result: 61.5

DIVISION BY 0

To avoid this error, make sure the denominator is not 0. You may use the `NULLIF()` function to replace it with a NULL, which results in a NULL for the entire expression:
count / NULLIF(count_all, 0)

JOIN

JOIN is used to fetch data from multiple tables. To get the names of products purchased in each order, use:
SELECT
`orders.order_date,`
`product.name AS product,`
`amount`
FROM `orders`
JOIN `product`
ON `product_id = orders.product_id;`

[Learn more about JOINS in our interactive SQL JOINs course.](#)

INSERT

To insert data into a table, use the INSERT command:
INSERT INTO `category`
VALUES
`(1, 'Home and Kitchen'),`
`(2, 'Clothing and Apparel');`

You may specify the columns to which the data is added. The remaining columns are filled with predefined default values or NULLs.
INSERT INTO `category (name)`
VALUES `('Electronics');`

UPDATE

To update the data in a table, use the UPDATE command:
UPDATE `category`
SET
`is_active = true,`
`name = 'Office'`
WHERE `name = 'Office';`

DELETE

To delete data from a table, use the DELETE command:
DELETE FROM `category`
WHERE `name IS NULL;`

[Check out our interactive course \[How to INSERT, UPDATE, and DELETE data in SQL.\]\(#\)](#)

Database Schema

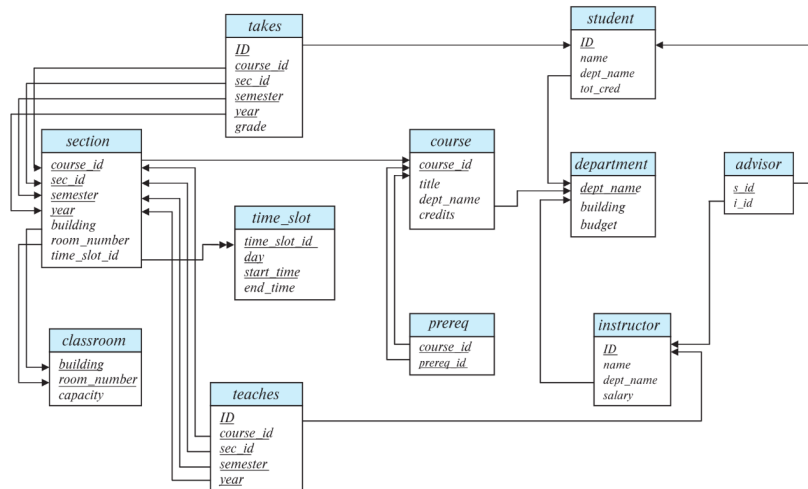
Here are the tables you'll find for the database used in the midterm. Your queries will be run against two versions of the database. One of the databases will be much smaller and only contain a subset of the information.

```
CREATE TABLE classroom
(
    building      varchar(15),
    room_number   varchar(7),
    capacity      numeric(4,0),
    primary key (building, room_number)
)
```

```
)
CREATE TABLE department
    (dept_name    varchar(20),
     building     varchar(15),
     budget       numeric(12,2) check (budget > 0),
     primary key (dept_name)
    )
CREATE TABLE course
    (course_id    varchar(8),
     title        varchar(50),
     dept_name    varchar(20),
     credits      numeric(2,0) check (credits > 0),
     primary key (course_id),
     foreign key (dept_name) references department (
on delete set null
    )
CREATE INDEX idx_course_dept_name ON Course(dept_name)
CREATE TABLE instructor
    (ID           varchar(5), -- instructor's ID
     name         varchar(20) not null,
     dept_name    varchar(20),
     salary       numeric(8,2) check (salary > 29000)
     primary key (ID),
     foreign key (dept_name) references department (
on delete set null
    )
CREATE INDEX idx_instructor_id ON Instructor(ID)
CREATE INDEX idx_instructor_dept_name ON Instructor(dept_name)
CREATE TABLE section
    (course_id    varchar(8),
     sec_id       varchar(8),
     semester     varchar(6)
     check (semester in ('Fall', 'Winter', 'Spring',
     year         numeric(4,0) check (year > 1701 and
     building     varchar(15),
     room_number  varchar(7),
     time_slot_id varchar(4),
     primary key (course_id, sec_id, semester, year)
     foreign key (course_id) references course (course_id)
on delete cascade,
```

```
foreign key (building, room_number) references
on delete set null
)
CREATE INDEX idx_section_year ON Section(year)
CREATE INDEX idx_section_semester ON Section(semester)
CREATE TABLE teaches
(
    ID          varchar(5), -- instructor's ID
    course_id   varchar(8),
    sec_id      varchar(8),
    semester    varchar(6),
    year        numeric(4,0),
    primary key (ID, course_id, sec_id, semester, year),
    foreign key (course_id, sec_id, semester, year)
    references section (course_id, sec_id, semester)
    on delete cascade,
    foreign key (ID) references instructor (ID)
    on delete cascade
)
CREATE INDEX idx_teaches_id ON Teaches(ID)
CREATE INDEX idx_teaches_year ON Teaches(year)
CREATE INDEX idx_teaches_semester ON Teaches(semester)
CREATE TABLE student
(
    ID          varchar(5),
    name        varchar(20) not null,
    dept_name   varchar(20),
    tot_cred    numeric(3,0) check (tot_cred >= 0),
    primary key (ID),
    foreign key (dept_name) references department (dept_name)
    on delete set null
)
CREATE UNIQUE INDEX idx_student_id ON Student(ID)
CREATE INDEX idx_student_dept_name ON Student(dept_name)
CREATE TABLE takes
(
    ID          varchar(5), -- Student ID
    course_id   varchar(8),
    sec_id      varchar(8),
    semester    varchar(6),
    year        numeric(4,0),
    grade       varchar(2),
    primary key (ID, course_id, sec_id, semester, year)
```

```
foreign key (course_id, sec_id, semester, year)
references section (course_id, sec_id, semester
on delete cascade,
foreign key (ID) references student (ID)
on delete cascade
)
CREATE INDEX idx_takes_id ON Takes(ID)
CREATE INDEX idx_takes_semester ON Takes(semester)
CREATE INDEX idx_takes_year ON Takes(year)
CREATE TABLE advisor
(s_ID      varchar(5),
i_ID      varchar(5),
primary key (s_ID),
foreign key (i_ID) references instructor (ID)
on delete set null,
foreign key (s_ID) references student (ID)
on delete cascade
)
CREATE INDEX idx_advisor_instructor_id ON Advisor(i_
CREATE TABLE time_slot
(time_slot_id  varchar(4),
day           varchar(1),
start_hr      numeric(2) check (start_hr >= 0
start_min     numeric(2) check (start_min >= 0
end_hr        numeric(2) check (end_hr >= 0 an
end_min       numeric(2) check (end_min >= 0 a
primary key (time_slot_id, day, start_hr, start
)
CREATE TABLE prereq
(course_id    varchar(8),
prereq_id    varchar(8),
primary key (course_id, prereq_id),
foreign key (course_id) references course (cour
on delete cascade,
foreign key (prereq_id) references course (cour
)
```

Scratch area

The following scratch space can be used to help develop and test queries against a database described above. The database used by the exam grader will be different.

Execute
Minimize Output

Questions For a total of 100 points

SQL Queries 65 points

In this section, you will write SQL queries for the university described in the book. Your queries will be tested immediately against two different databases. If your queries output matches the expected output, the

displayed answers will be outlined in green. Your actual score will be determined when your query is tested against a different database but green feedback should mean that you are on track to receive full credit.

List.All.1: Create a SQL query for the relational expression:

$\Pi_{\text{name}} (\sigma_{\text{dept_name} = \text{'Comp. Sci.'}} (\text{student}))$

Execute
4 points
Minimize Output

Relational.Algebra.1: Create a SQL query for the relational expression:

$\Pi_{\text{title}} (\sigma_{\text{semester} = \text{'Fall'} \wedge \text{dept_name} = \text{'Astronomy'}} (\rho_S(\text{Section}) \bowtie_{\text{C.course_id} = \text{S.course_id}} \rho_C(\text{Course})))$

$\Pi_{\text{title}} (\sigma_{\text{semester} = \text{'Winter'} \wedge \text{dept_name} = \text{'Astronomy'}} (\rho_S(\text{Section}) \bowtie_{\text{C.course_id} = \text{S.course_id}} \rho_C(\text{Course})))$

$\Pi_{\text{title}} (\sigma_{\text{semester} = \text{'Spring'} \wedge \text{dept_name} = \text{'Astronomy'}} (\rho_S(\text{Section}) \bowtie_{\text{C.course_id} = \text{S.course_id}} \rho_C(\text{Course})))$

$\Pi_{\text{title}} (\sigma_{\text{semester} = \text{'Summer'} \wedge \text{dept_name} = \text{'Astronomy'}} (\rho_S(\text{Section}) \bowtie_{\text{C.course_id} = \text{S.course_id}} \rho_C(\text{Course})))$

Execute 4 points Minimize Output

Relational.Algebra.In.Words.1: Which of the following best represents in words what **Relational.Algebra.1** is asking for?

- A List the courses that are offered by the Astronomy department Fall semester
- B List the unique course titles that are offered by the Astronomy department all semesters
- C List the unique course titles that are offered by the Astronomy department Fall semester but not at other times
- D List the courses that are offered by the Astronomy department all semesters
- E List the unique course titles that are offered by the Astronomy department Fall but not Spring, Summer semester(s)
- F List the unique course titles that are offered by the Astronomy department Fall but not Winter semester(s)
- H All of the above
- I None of the above

▼ 4 points

Did.Not.Take.1: List the name and total number of credits for students in the Comp. Sci. department who did not take the comp. sci. course named 'Game Design'.

Order the list in descending order by the name.

Execute 9 points Minimize Output

How.Many.Students.1: Find the number of students in the Comp. Sci. department.

Execute 4 points Minimize Output

New.Instructor.1: Add a new instructor named 'Pauli Murray', ID 'pauli', with a salary of 75,000 in the department of History.

NOTE₁: You may be assured that no other instructor has the ID pauli and so you may hardcode your query using it.

NOTE₂: With all SQL statements like this that modify the database, you either need to refresh the web page before rerunning or handle the fact that the statement was previously run.

Execute 9 points Minimize Output

New.Advisor.1: Make the new instructor from the previous question, Pauli Murray, the advisor to student 'Bates'.

NOTE₁: You need to get the previous question correct to get this question correct.

NOTE₂: You can hardcode the advisor id of pauli, but you **MUST** programmatically determine student Bates's ID. In other words, do not print out Bates's ID and then hardcode it. Programmatically determine it within the SQL code!

NOTE₃: You can be assured that the student name Bates is unique.

NOTE₄: With all SQL statements like this that modify the database, you either need to refresh the web page before rerunning or handle the fact that the statement was previously run.

Execute 9 points Minimize Output

Strings.1: Select all the student names who have a name with second character of a.

Order the result alphabetically by the student's department name. For all students within each department sort by the student's name.

Execute 4 points Minimize Output

Instructor.Count.1: List the names of Cybernetics instructors and the number of times they have taught a course in Taylor building. If they have never taught a course in that building, express the count as Null

Execute 9 points Minimize Output

Student.Takes.All.1: List the names of all the students who have taken all the courses in the Cybernetics department.

Execute
9 points
Minimize Output

Joins 20 points

In this section, you will show your knowledge of the different join types.

Relation instance A. Attribute names are in uppercase, **bold**, and **blue** in the first row. Tuples are in the rows that follow.

A1	A2	A3	A4
a11	a21	a31	a41
a12	a22	a32	a42
a13	a23	a33	a43

Relation instance B. Attribute names are in uppercase, **bold**, and **blue** in the first row. Tuples are in the rows that follow.

B1	B2	B3	B4	B5	B6
b11	b21	b31	a41	b51	b61
b12	b22	b32	a42	b52	b62

Cartesian.Product.1: Given the relational instances above, what is the result of the following join?

```
SELECT DISTINCT A2, B2, B5
FROM A, B
ORDER BY A2, B2, B5
```

NOTE: If after completing the rows returned from the query and there are rows remaining in the table, entry ROWUNUSED for the unused rows.

A2	B2	B5
a21 ▼	b21 ▼	b51 ▼
a21 ▼	b22 ▼	b52 ▼
a22 ▼	b21 ▼	b51 ▼
a22 ▼	b22 ▼	b52 ▼
a23 ▼	b21 ▼	b51 ▼
a23 ▼	b22 ▼	b52 ▼

5 points

Inner.Join.1: Given the relational instances above, what is the result of the following join?

```
SELECT DISTINCT A2, B2, B5
FROM A, B
WHERE A4 = B4
ORDER BY A2, B2, B5
```

NOTE: If after completing the rows returned from the query and there are rows remaining in the table, entry ROWUNUSED for the unused rows.

A2	B2	B5
a21 ▼	b21 ▼	b51 ▼
a22 ▼	b22 ▼	b52 ▼
ROWUNUSED ▼	ROWUNUSED ▼	ROWUNUSED ▼
ROWUNUSED ▼	ROWUNUSED ▼	ROWUNUSED ▼
ROWUNUSED ▼	ROWUNUSED ▼	ROWUNUSED ▼
ROWUNUSED ▼	ROWUNUSED ▼	ROWUNUSED ▼

5 points

Left.Join.1: Given the relational instances above, what is the result of the following join?


```

SELECT DISTINCT A2, B2, B5
FROM A LEFT JOIN B
      ON A4 = B4
ORDER BY A2, B2, B5

```

NOTE: If after completing the rows returned from the query and there are rows remaining in the table, entry ROWUNUSED for the unused rows.

A2	B2	B5
<input type="text" value="a21"/>	<input type="text" value="b21"/>	<input type="text" value="b51"/>
<input type="text" value="a22"/>	<input type="text" value="b22"/>	<input type="text" value="b52"/>
<input type="text" value="a23"/>	<input type="text" value="NULL"/>	<input type="text" value="NULL"/>
<input type="text" value="ROWUNUSED"/>	<input type="text" value="ROWUNUSED"/>	<input type="text" value="ROWUNUSED"/>
<input type="text" value="ROWUNUSED"/>	<input type="text" value="ROWUNUSED"/>	<input type="text" value="ROWUNUSED"/>
<input type="text" value="ROWUNUSED"/>	<input type="text" value="ROWUNUSED"/>	<input type="text" value="ROWUNUSED"/>

5 points

Full.Outer.Join.1: Given the relational instances above, what is the result of the following join?

```

SELECT DISTINCT A2, B2, B5
FROM A FULL OUTER JOIN B
      ON A4 = B4
ORDER BY A2, B2, B5

```

NOTE: If after completing the rows returned from the query and there are rows remaining in the table, entry ROWUNUSED for the unused rows.

A2	B2	B5
<input type="text" value="a21"/>	<input type="text" value="b21"/>	<input type="text" value="b51"/>
<input type="text" value="a22"/>	<input type="text" value="b22"/>	<input type="text" value="b52"/>
<input type="text" value="a23"/>	<input type="text" value="NULL"/>	<input type="text" value="NULL"/>
<input type="text" value="ROWUNUSED"/>	<input type="text" value="ROWUNUSED"/>	<input type="text" value="ROWUNUSED"/>
<input type="text" value="ROWUNUSED"/>	<input type="text" value="ROWUNUSED"/>	<input type="text" value="ROWUNUSED"/>
<input type="text" value="ROWUNUSED"/>	<input type="text" value="ROWUNUSED"/>	<input type="text" value="ROWUNUSED"/>

5 points

Chapter Reading Review 15 points

Databases.Are.1: Select the letter of the statement below that is false.

- A Database systems are designed to store large bodies of information
- B The relational data model is the only model used for storing data in databases
- C The management of data involves both the definition of structures for the storage of information and the provision of mechanisms for the manipulation of information
- D The database system must provide for the safety of the information stored in the face of system crashes or attempts at unauthorized access
- E If data is shared among users, the database system must avoid possible anomalous results

▼ 1 point

Data.Manipulation.Language.1: Select the letter of the data manipulation language statement below that is false or select E if none of them are false.

- A The data manipulation language is used to specify the database schema.
- B The data manipulation language enables users to access data.
- C The INSERT statement is a data manipulation language statement.
- D The data manipulation language is widely used in retrieving data from relational databases.
- E None of the above are false.

▼ 1 point

For the following fill-in-the-blank-questions match the letter for the phrase that best answers the following questions.

- A A collection of iterrelated data and the programs to access the data
- B Edgar "Ted" Codd
- C Data documentation language
- D A collection of data
- E Standard query language
- F Fred Brooks
- G Data manipulation language
- H None of these

Word.Definition.1: Letter representing the best answer to the question:

Question	Letter of definition
----------	----------------------

What is a database?	<input type="text" value="v"/>
What is a database management system?	<input type="text" value="v"/>
Who proposed and later won a Turing award for the relational database model?	<input type="text" value="v"/>
What is an example of a database language?	<input type="text" value="v"/>

4 points


Each of the items in below describe a problem solved by a database. The table contains terms of features of a database. For each database feature, select the letter of the problem that it solves.

- A Larger amounts of storage and the cost to access data in various formats
- B Confusion of physical versus logical layer
- C Efficient and convenient access to data
- D All portions or no portions of a logical unit of work completes. The all-or-none requirement
- E Physical layer data abstraction
- F University payroll personnel should see financial information but not academic or grade records
- G Multiple data updates to a field that occur near simultaneously to each updater without inconsistent results
- H None of these

DB.Solves.1 For each database feature below, select the letter of the problem that it solves.

Database feature	Problem solved
Atomicity	<input type="text" value="v"/>
Security	<input type="text" value="v"/>
Concurrency	<input type="text" value="v"/>

4 points



Following are relational algebra statements:

Note: During the test, this question incorrectly had σ where it should have had Π .

$$\mathbf{A} \Pi_{\text{course_id, title}}(\textit{Teaches} \bowtie \textit{Course})$$
$$\mathbf{B}(\Pi_{\text{course_id, title}} \textit{Teaches}) \bowtie (\Pi_{\text{dept_name, title}} \textit{Course})$$
$$\mathbf{c} \Pi_{\text{course_id, title}}(Teaches \bowtie_{Teaches.course_id=Course.course_id} Course)$$
$$\mathbf{D} \Pi_{\text{course_id, title}}((\Pi_{\text{course_id, year}} \textit{Teaches}) \bowtie \textit{Teaches.course_id} = \textit{Course.course_id} \textit{Course})$$
$$\mathbf{E} \Pi_{\text{course_id, title}}(Teaches \bowtie_{Teaches.course_id=Course.course_id} (\Pi_{\text{course_id, title}} Course))$$

F All of the above relational algebra produces the same results

G There are more than two different results


Teachers.Teach.1: Which of the above relational algebra statements produce different results than all the others?

Select **F** or **G**, respectively, if they all produce the same results or more than two different results.

▼ 3 points

Relational.Algebra.Operation.1: What relational algebra operation is represented by the Greek letter sigma, σ ?

1 point



Reading.Schema.1: Given the schema diagram above, can a department be in two different buildings?

▼ 1 point

Pledge your worksheet

DO NOT PUT YOUR NAME ON THE WORKSHEET UNTIL YOU ARE READY TO SUBMIT IT.

You will be given grace on any two assignments that you want to resubmit. But you should consider that only the first submission of a worksheet or any assignment will be graded.

UNC Honor Pledge: I certify that no unauthorized assistance has been received or given in the completion of this work. This unauthorized assistance includes, but is not limited to, copying of another student's answers.

Enter your full name

If other students helped in the completion of this worksheet, give attribution to them by entering their Onyen below. Help can be in the form of a work group, online conversation, etc. It must **not** include copying answers. If no other student helped with the completion of this worksheet, **None** should be entered.

Enter the Onyen of any students who assisted in this worksheet's co

Done !

Submit your work

Did you **pledge your work** above?

Did you **acknowledge any student assistance** above?

Submit